

AMAZON SPONSORED PRODUCTS

Analysis Report

Reporting Period: July 1 – July 30, 2026

Scope: 2,435 campaigns · 61 portfolios · United States Marketplace (Amazon, USD)

By : Ayush Aditya

Data sources: Advertised Product, Search Term & Targeting reports as Provided by California Design

Introduction

Over the 30-day period analyzed, the Sponsored Products program generated \$5,445,552 in attributed sales from \$1,479,163 in ad spend — a blended ACOS of 27.2% and a return on ad spend (ROAS) of 3.68x, built on 134.9 million impressions and 620,891 clicks across 2,435 campaigns in 61 portfolios. Performance was stable rather than volatile: daily ACOS held in a 22–35% band across the month with no sustained upward or downward trend, indicating a mature, steady-state program rather than one in a growth or decline phase.

The account-level average masks meaningful dispersion underneath it. Five analytical lenses were applied — account trends, portfolio/campaign concentration, search-term efficiency, targeting-type mix, and top-of-search placement — and each surfaced a concrete, dollar-quantified opportunity:

- **Search-term waste:** 150,009 of 181,179 unique customer search terms (83%) produced zero orders, costing \$413,248 — 28% of total spend — with \$165,908 of that concentrated in a set of roughly 17,500 terms that are probably better removed.
- **Targeting-mix inefficiency:** Keyword targeting absorbs 63% of spend (\$936,418) yet posts the weakest ACOS (29.3%) of any targeting type — Amazon's own Auto-targeting and ASIN-targeting logic is outperforming manual keyword curation by 4–7 ACOS points.
- **Placement headroom:** 90% of spend sits in targets winning only 0–10% top-of-search impression share, where conversion rate is 15.4% — versus 26.8% for targets winning 50%+ of top-of-search, at a comparable ACOS.
- **Clear pause and scale lists:** 9 campaigns (\$30,421 spend, 58% ACOS) are underwater; 16 campaigns (\$31,277 spend, 9.72x ROAS) are proven and under-funded relative to their efficiency.

Collectively, the actionable waste and reallocation opportunities identified below touch roughly \$200,000 of the account's monthly spend — about 14% of the total budget — without requiring any change to overall program size. In Section 6 you will find recommendations to convert these findings into an action plan.


1. Account-Level Performance & Trends

The table below summarizes headline account metrics for the full period, followed by the daily trend in spend, sales and ACOS.

Metric	Value	Metric	Value
Total Spend	\$1,479,163	Total Orders	97,644
Total Sales (7-Day)	\$5,445,552	Total Units Sold	110, 411
ACOS	27.2%	Click-Through Rate	0.46%
ROAS	3.68x	Conversion Rate (Orders/Click)	15.7%
Impressions	134,940,036	Avg. Cost-Per-Click	\$2.38
Clicks	620,891	Active Campaigns / Portfolios	2,435 / 61

Table 1. Account-level KPIs, July 1–30, 2026.

Figure 1. Daily spend and sales (bars, left axis) against daily ACOS (line, right axis). ACOS ranged 22–35% with a mild dip mid-month (Jul 14–17) coinciding with an increase in spending increase, and no clear weekly seasonality is visible at this aggregation.




Two observations stand out. First, sales tracked spend closely all month (correlation is visually tight in Figure 1), meaning the account was not spend-constrained — additional budget on any given day would very likely have converted at a similar rate, which supports the scale recommendations in Section 6. Second, ACOS variance across the month (a ~13-point swing between the best and worst day) is wider than typical day-of-week seasonality would explain on its own, suggesting that a portion of the daily ACOS movement is driven by which campaigns happened to spend more that day rather than by demand shifts — a hypothesis Section 2’s campaign-level data supports.

2. Portfolio & Campaign Concentration

Spend is concentrated in a small number of portfolios and campaigns, which is normal for a mature account — but the efficiency of that concentration determines whether it's healthy or risky.

Figure 2. Top 15 portfolios by spend, plotted against ACOS; bubble size is total sales. Portfolios above the dashed line run less efficiently than the account average (27.2% ACOS).



Portfolio_010, Portfolio_020 and Portfolio_050 are the three largest portfolios by spend (a combined \$398,563, or 27% of total budget) and all sit close to the account-average ACOS band (28–31%) — large, stable, unremarkable performers. Portfolio_022 and Portfolio_043 are similarly large but run 4–5 ACOS points hotter than average, worth a closer bid audit. The standout is Portfolio_039: at \$54,570 in spend it is mid-sized, but its 13.5% ACOS and 7.42x ROAS make it the most efficient large portfolio in the account by a wide margin — a strong candidate for a deliberate budget increase (see Section 6).

Portfolio	Spend	Sales	ACOS	ROAS
Portfolio_010	\$140,649	\$453,826	31.0%	3.23x
Portfolio_020	\$137,150	\$494,180	27.8%	3.60x
Portfolio_050	\$120,764	\$398,908	30.3%	3.30x
Portfolio_022	\$110,524	\$336,357	32.9%	3.04x
Portfolio_043	\$86,803	\$274,848	31.6%	3.17x
Portfolio_039 (efficiency standout)	\$54,570	\$404,838	13.5%	7.42x

Table 2. Top-spending portfolios versus the efficiency standout, Portfolio_039.

Campaign-level concentration




Figure 3. Pareto curve: campaigns ranked by spend (x-axis) against their cumulative share of total spend and total sales (y-axis).

The top 20% of campaigns by spend (487 of 2,435) account for 90.2% of spend and 89.8% of sales — the two curves track almost on top of each other. That near-equality is the healthy version of concentration: the account isn't over-investing in a head of inefficient campaigns while the long tail quietly converts, nor the reverse. The long tail itself is low-risk: 910 campaigns (37% of all campaigns) produced zero sales, but together they spent only \$8,248 (0.6% of total budget) — evidence of disciplined campaign-level budget caps rather than runaway test spend.

The concentration risk that does exist is smaller and sits at the extremes: 9 campaigns spending over \$1,000 each carry an ACOS above 50% (a blended 58% ACOS on \$30,421), while 16 campaigns of similar spend run below 15% ACOS at a blended 9.72x ROAS. Both lists are compact enough to impact the action directly — see Section 6.

3. Search-Term Efficiency & Wasted Spend

The search-term report is the most granular lens available and reveals the account's largest single source of avoidable spend. Of 181,179 unique customer search terms that triggered an ad, 150,009 (82.8%) produced zero orders, together absorbing \$413,248 — 28.0% of total spend.




Figure 4. The twelve individual search terms with the highest spend and zero resulting orders.

Not all of this \$413,248 is equally actionable, so it's useful to split it by how concentrated the waste is per term:

Tier	Criteria	# Terms	Spend	Recommended Action
1	≥ \$20 spend, 0 orders	1,018	\$32,399	Add as exact negative keyword individually
2	\$5–20 spend, 0 orders	16,537	\$133,509	Batch-review by theme; negate at ad-group level
3	< \$5 spend, 0 orders (long tail)	132,454	\$247,341	Address via category/auto-negation rules, not manually

Table 3. Zero-order search terms segmented by tractability. Tiers 1 and 2 together (\$165,908) are the near-term actionable target.

On the other side of the ledger, a small number of search terms are extremely efficient and worth deliberately harvesting into their own exact-match campaigns if not already isolated: “california design den sheets” converts at an 11.5% ACOS on \$12,171 spend, and the near-branded “california design den pillow cases” converts at just 4.75% ACOS (21.0x ROAS) — the single most efficient term of meaningful volume in the account. Generic head terms like “100% cotton sheets” and “cotton sheets” still convert profitably (35–38% ACOS) but at roughly triple the ACOS of the branded terms, which is the expected pattern for unbranded, high-competition queries.

Match-type performance

Match Type	Spend	Share of Spend	ACOS	ROAS
Auto (no match type)	\$623,273	42.2%	24.8%	4.02x
Phrase	\$465,284	31.5%	29.3%	3.41x
Exact	\$321,166	21.7%	31.2%	3.21x
Broad	\$66,791	4.5%	22.0%	4.55x


Table 4. Spend and efficiency by keyword match type, from the search-term report.

This ordering is a genuine anomaly worth flagging rather than smoothing over: conventional Amazon PPC practice expects Exact match to be the most efficient tier, since it targets the tightest, highest-intent queries. Here Exact posts the worst ACOS (31.2%) of the four types, while Broad — typically the riskiest, most-waste-prone tier — posts the best (22.0%), albeit on a much smaller base of spend. A plausible explanation is that Exact-match keywords in this account were seeded broadly (including generic, non-purchase-intent queries) rather than curated from proven converters, while Broad match here is a small, well-curated set. This is explored further against the targeting report in Section 4.

4. Targeting Strategy: Manual vs. Automatic

The targeting report classifies every target into six groups: manual keyword targeting, four Auto-targeting sub-types (close-match, loose-match, substitutes, complements), and manual product/ASIN targeting.

Figure 5. Spend (bars, left axis) versus ACOS (bars, right axis) by targeting type.



Targeting Type	Spend	Share	ACOS	ROAS
Keyword Targeting (manual)	\$936,418	63.3%	29.3%	3.42x
Auto – Close-match	\$197,384	13.3%	24.8%	4.03x
Auto – Loose-match	\$123,192	8.3%	23.9%	4.19x
Auto – Substitutes	\$107,759	7.3%	23.7%	4.21x
Product (ASIN) Targeting	\$83,328	5.6%	24.6%	4.07x
Auto – Complements	\$31,082	2.1%	21.8%	4.59x

Table 5. Spend and efficiency by targeting type, from the targeting report.

The pattern here reinforces Section 3's match-type finding: every Auto-targeting sub-type and Product (ASIN) targeting outperforms manual Keyword targeting on ACOS, by 4.5 to 7.5 percentage points, despite Keyword targeting receiving nearly two-thirds of total budget. Amazon's own algorithmic matching (Auto) is finding more efficient traffic than the account's manually curated keyword list is. This doesn't mean keyword targeting should be abandoned — it typically offers more control over bidding and message-to-query matching once curated — but it does suggest the keyword list carries more low-intent, low-converting terms than the Auto and ASIN buckets do, consistent with the Exact-match anomaly noted in Section 3.

Top-of-search placement and conversion




Figure 6. Spend (bars, left axis) alongside conversion rate and ACOS (lines, right axis), bucketed by the share of top-of-search impressions each target won.

Conversion rate rises steadily and substantially with top-of-search impression share — from 15.4% for targets that rarely win the top slot (0–10% share) to 26.8% for targets that win it more than half the time — while ACOS does not meaningfully worsen across the same range (27.7% down to 26.0%). The account's problem is not that top-of-search placement is unprofitable; it's that 90% of spend sits in the 0–10% bucket, meaning the vast majority of targets are simply not winning that placement often enough to benefit from it. This points to a bid-strategy lever, not a targeting lever: raising bids specifically on already-converting targets to compete more aggressively for top-of-search real estate is likely to lift conversion rate without a proportional ACOS penalty, based on the pattern observed here.

5. Synthesis: What's Driving the 27.2% Account ACOS

Reading Sections 1–4 together, three structural factors explain most of the gap between the account's blended 27.2% ACOS and the far more efficient pockets found throughout the data (13.5% at the portfolio level, 4.8% at the search-term level, sub-25% across every non-keyword targeting type):

- **A long, mostly-immaterial tail of zero-converting search terms** (28% of spend, but heavily concentrated in tiny per-term amounts) dilutes otherwise-strong campaign performance.
- **Manual keyword curation lags Amazon's own algorithmic Auto-targeting** by several ACOS points, despite carrying the majority of the budget — the account is over-weighted toward its least efficient targeting method.
- **Bidding is not aggressive enough on proven winners** to capture top-of-search placement, leaving a conversion-rate uplift (up to ~11 points) largely unclaimed on 90% of spend.

None of these are demand problems — daily sales tracked spend tightly all month (Figure 1) and the campaign-level Pareto curve is healthy (Figure 3) — which is why the recommendations below are framed as reallocation and refinement, not budget cuts.

6. Recommendations & Budget Reallocation




Figure 7. Current spend sitting behind each recommendation below, sized by dollar exposure.

Priority 1 — Search-term negation triage (\$165,908 addressable)

- Add the 1,018 Tier-1 terms ($\geq \$20$ spend, 0 orders; \$32,399 total) as exact negative keywords this week — the highest-confidence, lowest-effort action in this report.
- Batch-review the 16,537 Tier-2 terms (\$5–\$20 spend; \$133,509 total) by theme (e.g. product-adjacent but non-purchase-intent phrases) and negate at the ad-group level rather than one by one.
- Leave Tier-3 (132,454 terms under \$5 each, \$247,341 total) alone at the individual level — the effort-to-recovery ratio doesn't justify manual negation; if tooling allows, apply category- or root-word-level automated negation rules instead.

Priority 2 — Rebalance the targeting mix

- Shift a portion of net-new budget growth toward Auto and ASIN targeting rather than expanding the Keyword-targeting bucket further, given its persistent 4–7 point ACOS disadvantage.
- Audit the Exact-match keyword list specifically — its 31.2% ACOS (the worst of any match type) suggests it was seeded with broader, lower-intent terms than Exact match is meant to carry; cross-reference against the Tier-1/2 negative-keyword list in Priority 1, since overlap is likely.

Priority 3 — Prune underperformers, fund proven winners

- Pause or substantially cut bids on the 9 campaigns running above 50% ACOS (\$30,421 spend, 58% blended ACOS, 1.72x ROAS) — reallocate that budget rather than continuing to fund it.
- Increase daily budgets and/or bids on the 16 campaigns running below 15% ACOS (\$31,277 spend, 9.72x blended ROAS); the flat daily spend-to-sales relationship in Figure 1 suggests these campaigns are not currently budget-capped in a way that would blunt the return on extra spend.
- Apply the same logic to Portfolio_039 specifically (13.5% ACOS, 7.42x ROAS at \$54,570 spend) — the single most efficient large portfolio in the account and a natural home for reallocated budget.

Priority 4 — Bid up for top-of-search share on proven converters

- For targets already converting profitably but sitting in the 0–10% top-of-search impression-share band (90% of spend), test bid increases specifically aimed at winning more top-of-search placements — Figure 6 shows conversion nearly doubles at higher placement-share tiers without a corresponding ACOS penalty.

Directional impact: reallocating the ~\$165,908 in Priority 1 waste at the account's blended 3.68x ROAS implies roughly \$610,000 in potential incremental sales — offered as a ceiling estimate based on the account average, not a guarantee, since reallocated spend will not automatically convert at the account mean. Combined, Priorities 1–3 touch approximately \$200,000 of monthly spend (≈14% of the total budget) through re-targeting rather than net budget growth, which is the more defensible first move before requesting additional program spend.